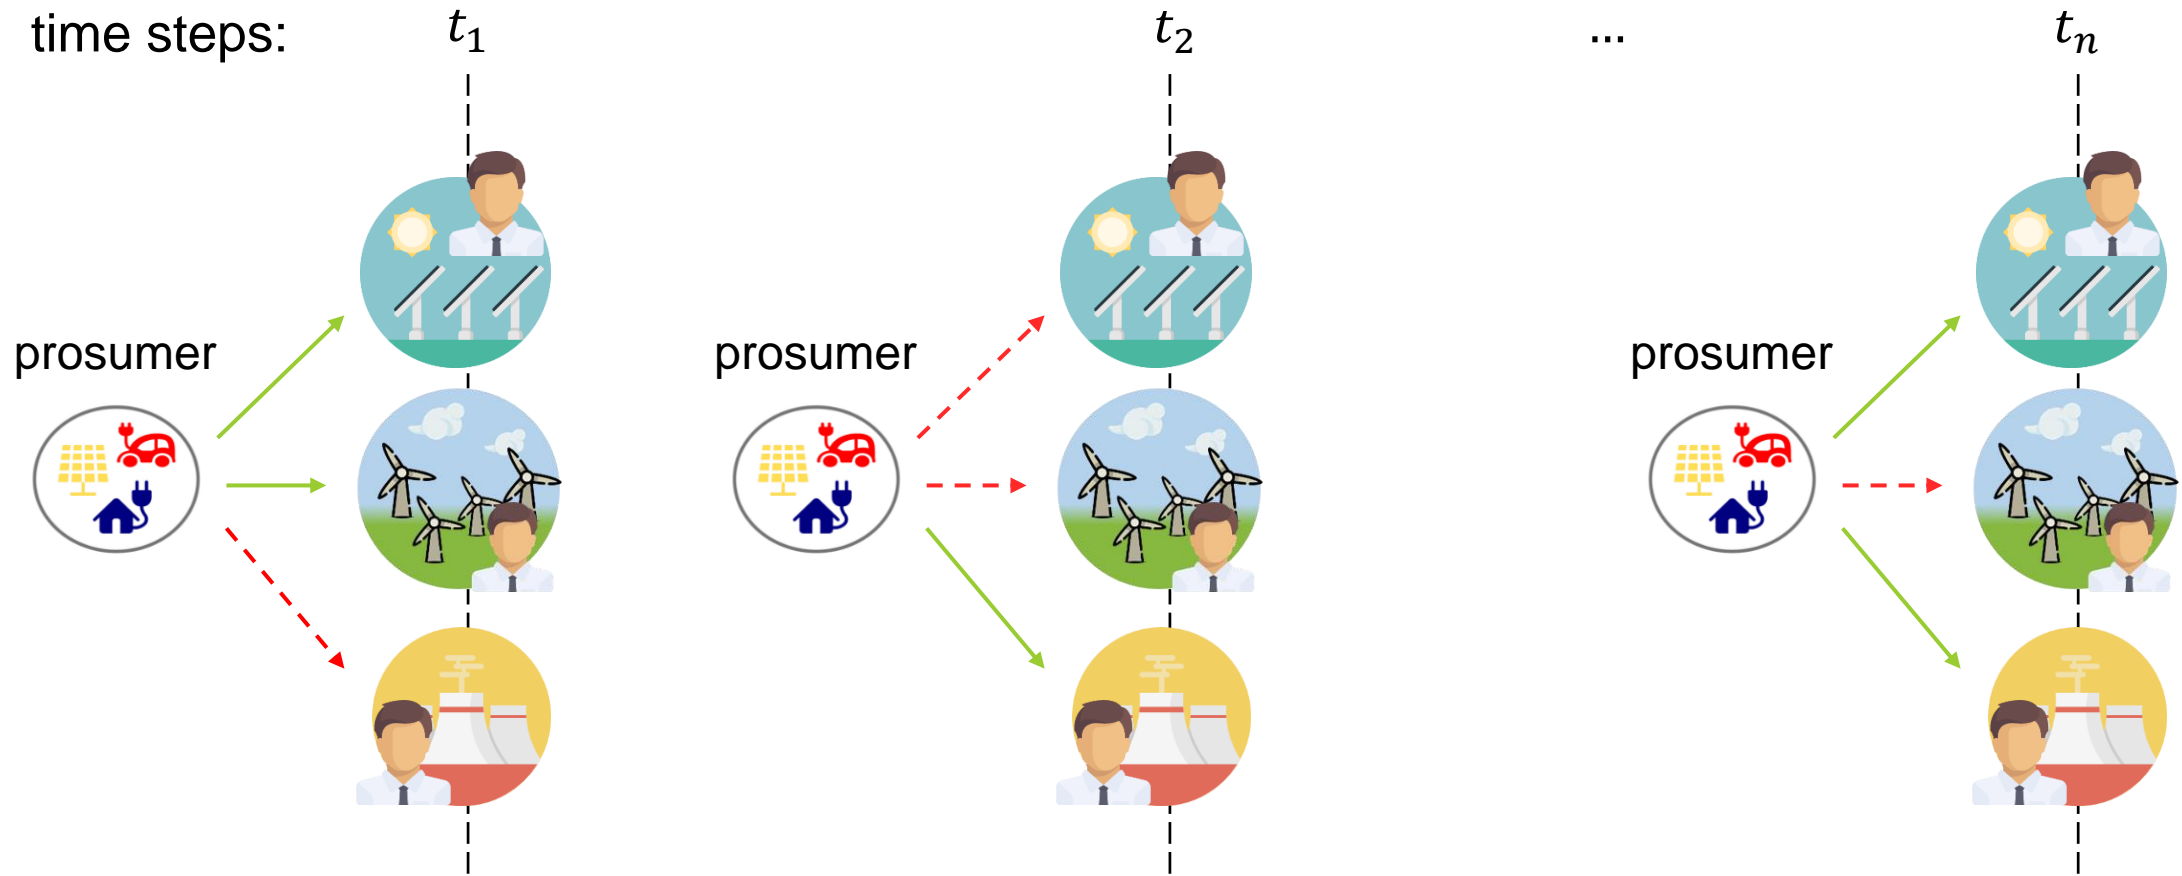




Tiago Sousa, Pierre Pinson

# Motivation

time steps:

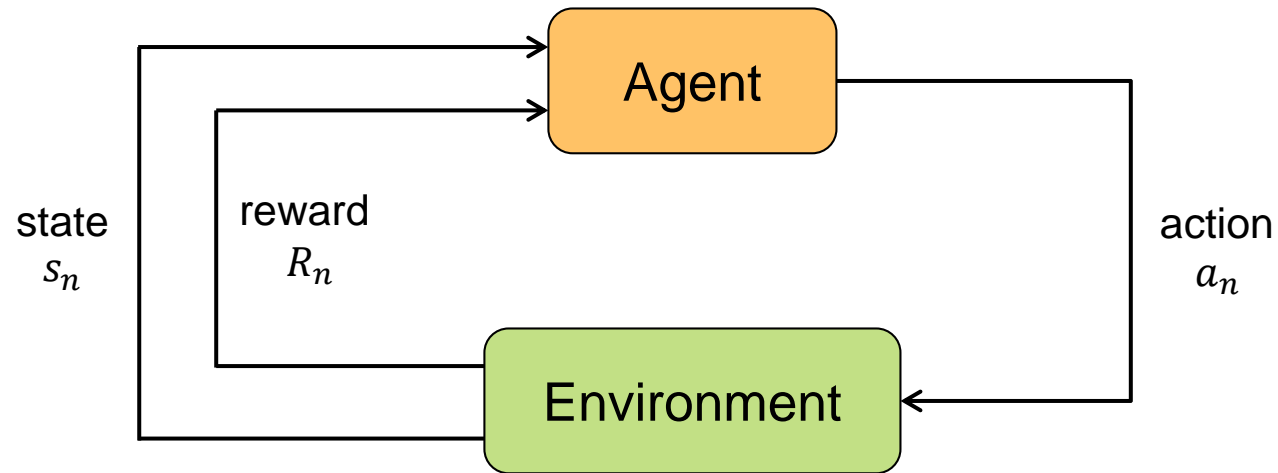


— Yes

— No

# Reinforcement learning

- We can design the P2P negotiation as Reinforcement Learning
  - Agent learns the optimal policy by interacting with the environment (P2P market)



# Outline

- Reinforcement Learning approach
  - Multi-Armed Bandit
- Test case and results
- Conclusions and next steps

# Multi-Armed Bandit

Agent

prosumer



<i>Step <math>n</math></i>	<i>Arm 1</i>	<i>Arm 2</i>	<i>Arm 3</i>
<b>1</b>	1	0	0
<b>2</b>	0	1	0
<b>3</b>	1	0	0
<b>4</b>	0	0	1

actions  $a_n$

Step  $n \neq$  time  $t$

# Multi-Armed Bandit

Agent

prosumer



Environment

<i>Step <math>n</math></i>	<i>Arm 1</i>	<i>Arm 2</i>	<i>Arm 3</i>	<i>Reward</i>
<b>1</b>	1	0	0	1
<b>2</b>	0	1	0	0
<b>3</b>	1	0	0	0
<b>4</b>	0	0	1	1

Total reward

$$R_{Total} = \mathbb{E}(R_n) = \frac{1}{N} \sum R_n$$

$$a_n^* = \operatorname{argmax} R_{Total}$$

# Multi-Armed Bandit

Agent

prosumer



Environment

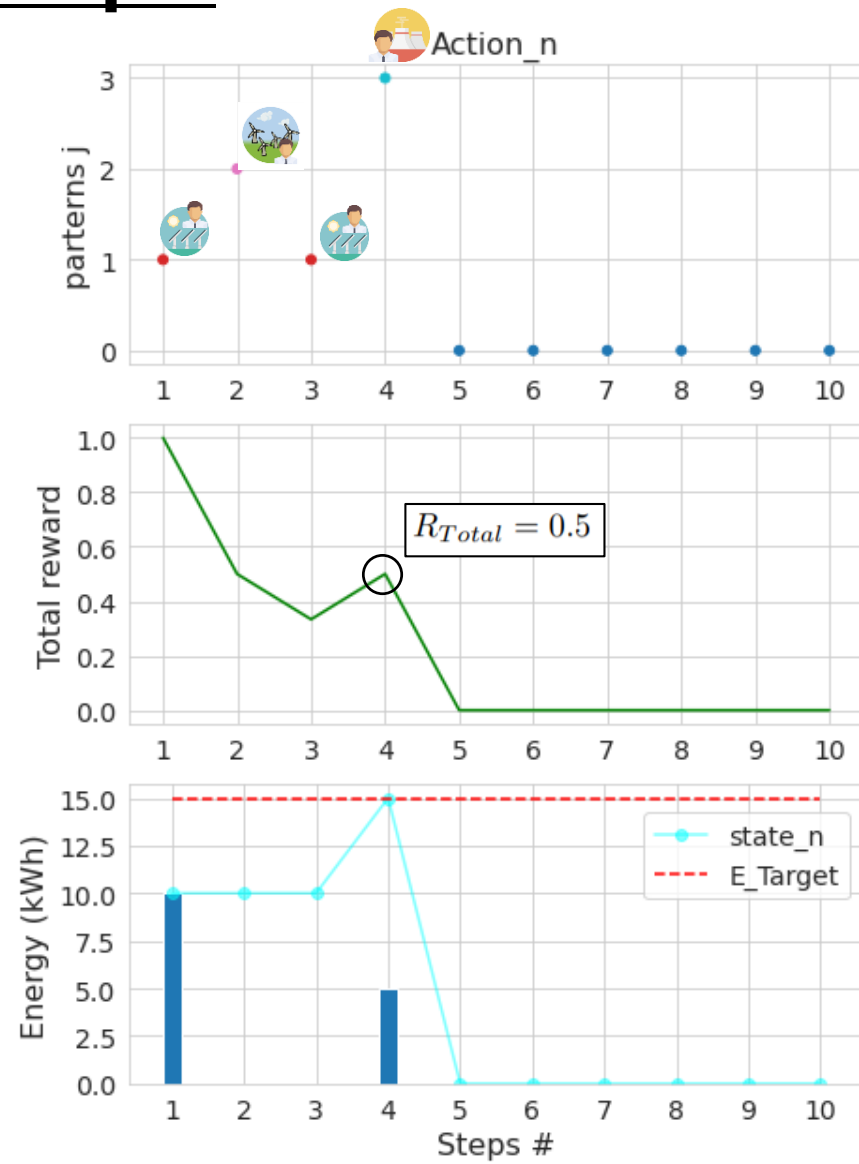
<i>Step <math>n</math></i>	<i>Arm 1</i>	<i>Arm 2</i>	<i>Arm 3</i>	<i>Reward</i>	<i>Energy (kWh)</i>
<b>1</b>	1	0	0	1	10
<b>2</b>	0	1	0	0	0
<b>3</b>	1	0	0	0	0
<b>4</b>	0	0	1	1	5

State per step  $n$

$$s_n = \sum E_n(j) R_n$$

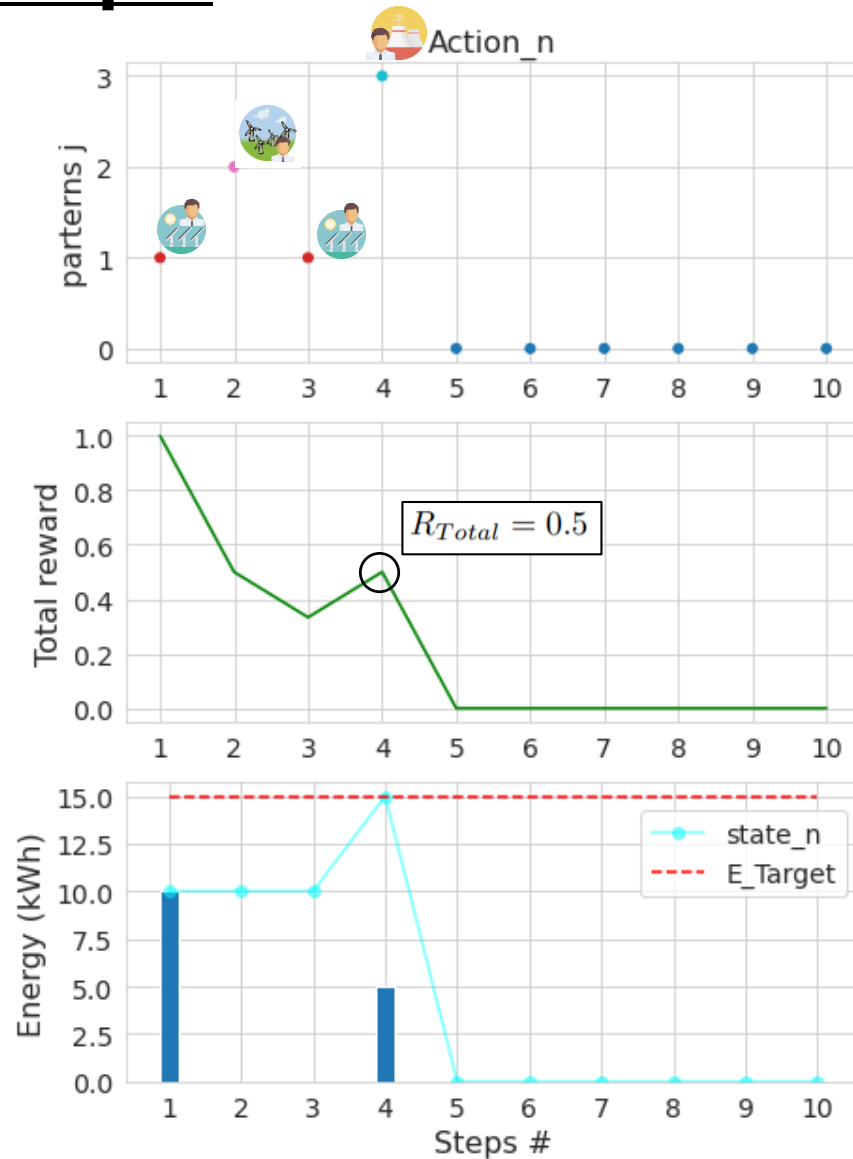
$$s_n \leq E_{Target} \quad \text{Stopping condition}$$

# Example





# Example





- This iterative process is an episode

- We terminate when

$$- s_n = E_{Target}$$

episode = time  $t$

# How to differentiate between episodes?



# Multi-Armed Bandit

- Reward is a random variable

## Bernoulli distribution

$$R_n(j) \sim \text{B}(1, p_j)$$

For a large number of steps  $n$ :

$$R_n(j) \approx p_j$$



<i>Step <math>n</math></i>	<i>Arm 1</i>
<b>1</b>	1
<b>2</b>	0
<b>3</b>	1
<b>4</b>	0

## Environment

<i>Reward</i>
1
-
0
-

# Multi-Armed Bandit


- Reward is a random variable

## Bernoulli distribution

$$R_n(j) \sim \text{B}(1, p_j)$$

For a large number of steps  $n$ :

$$R_n(j) \approx p_j$$



		Environment
Step $n$	Arm 1	Reward
1	1	1
2	0	-
3	1	0
4	0	-


## Action-Value function

Estimator of  $p_j$  for every arm  $j$

$$Q_n(j) = \frac{1}{N_j} \sum R_n(j) = \hat{p}_n(j)$$

# Multi-Armed Bandit


- Learning algorithms estimate the Action-Value function  $Q_n(j)$ 
  - Random
  - $\epsilon$ -greedy
  - Thompson Sampler
  - Upper Confidence Bound





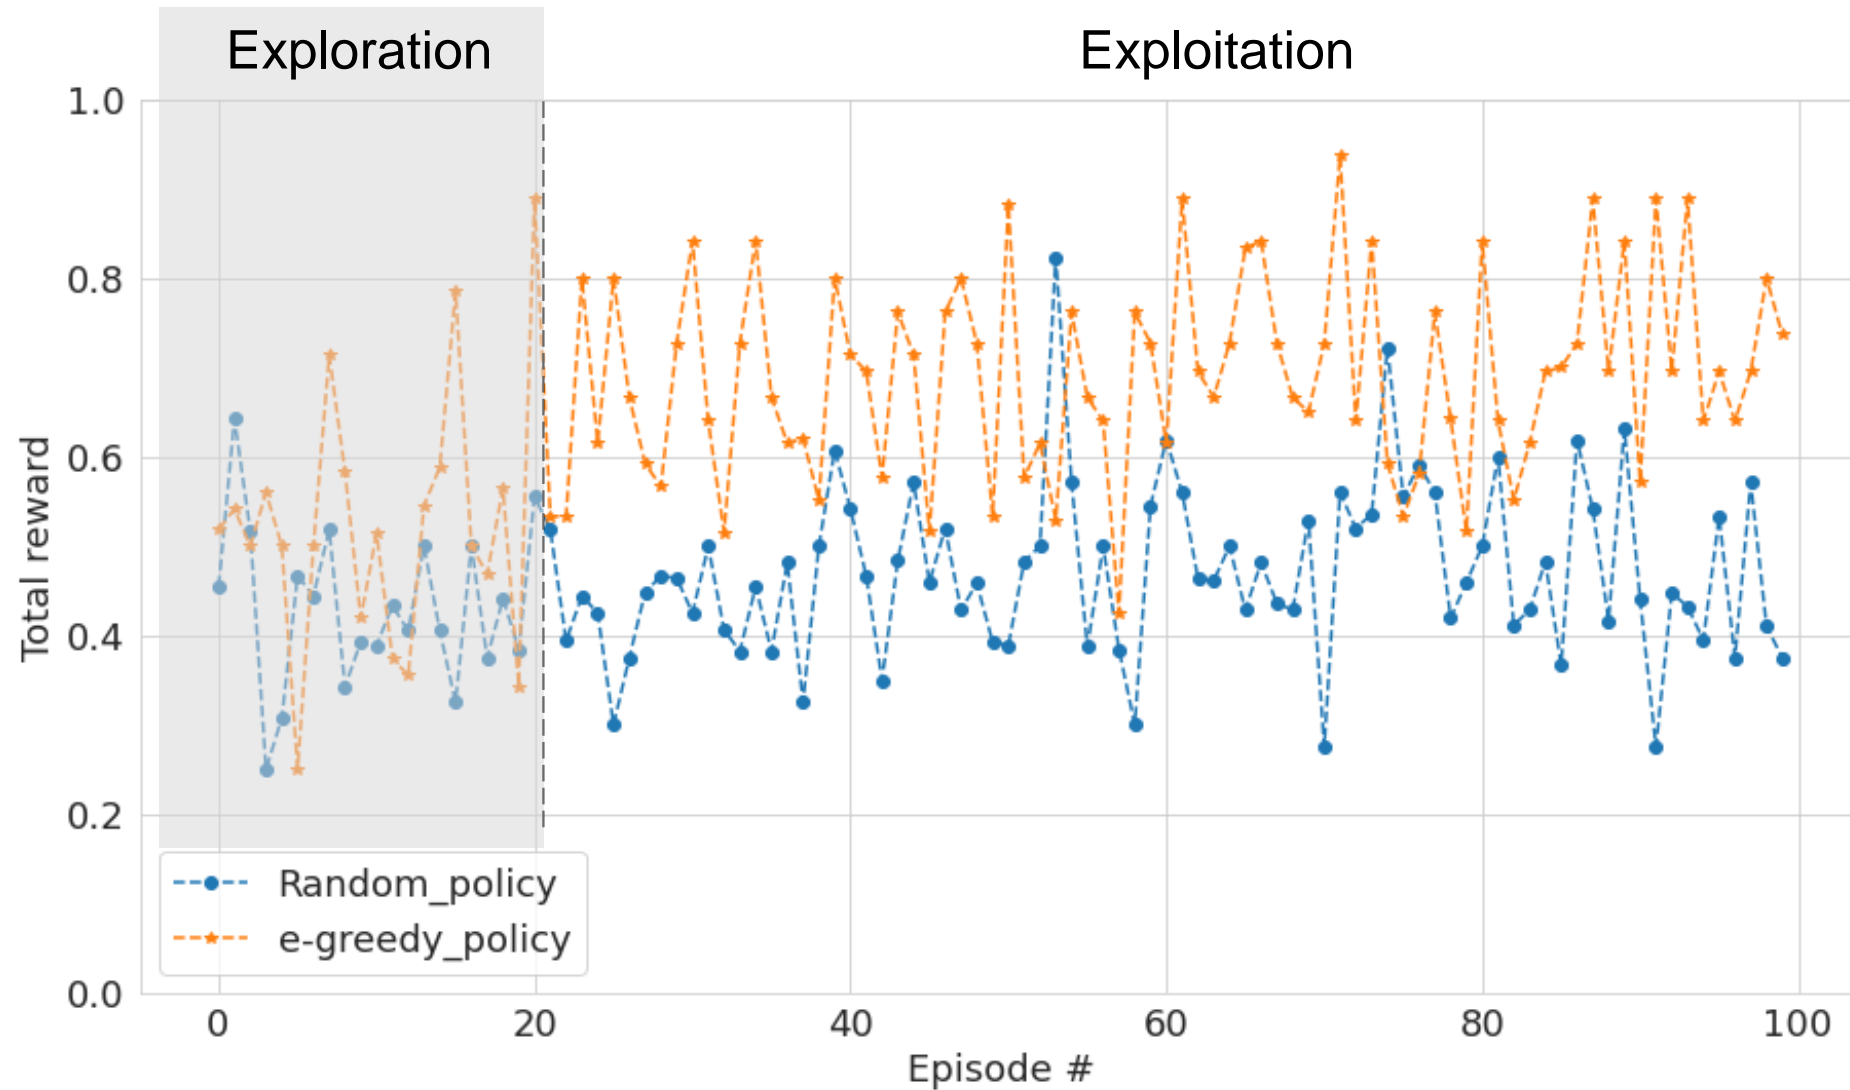
# Test case

- Case with 15 partners  $j$
- Test for 100 episodes
  - $E_{Target} = 15$  kWh
- Learning algorithms:
  - Random
  - $\epsilon$ -greedy



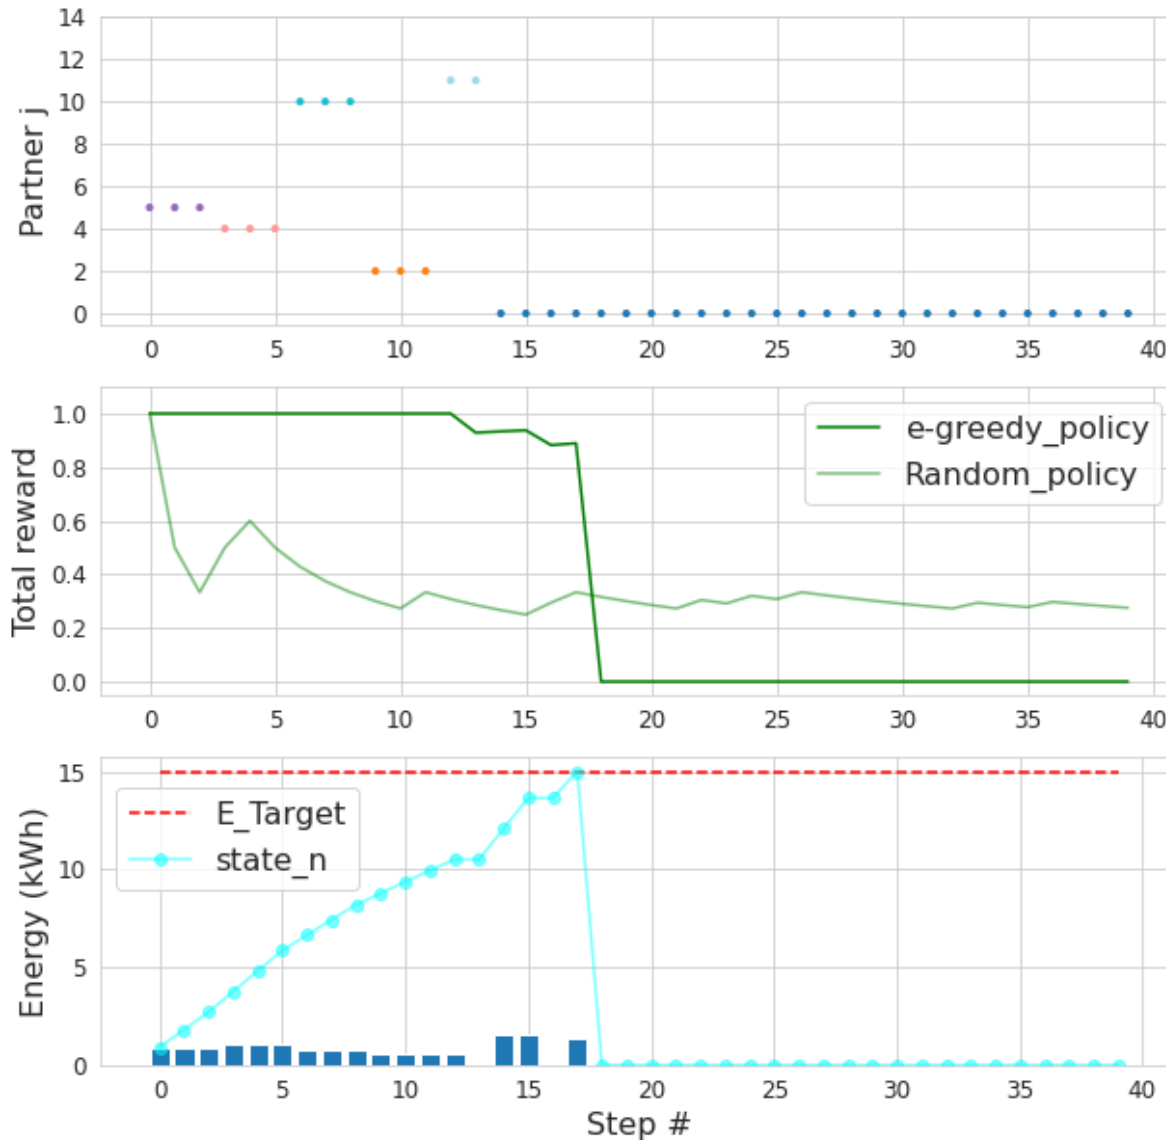


# Test case





# Epsilon-Greedy algorithm

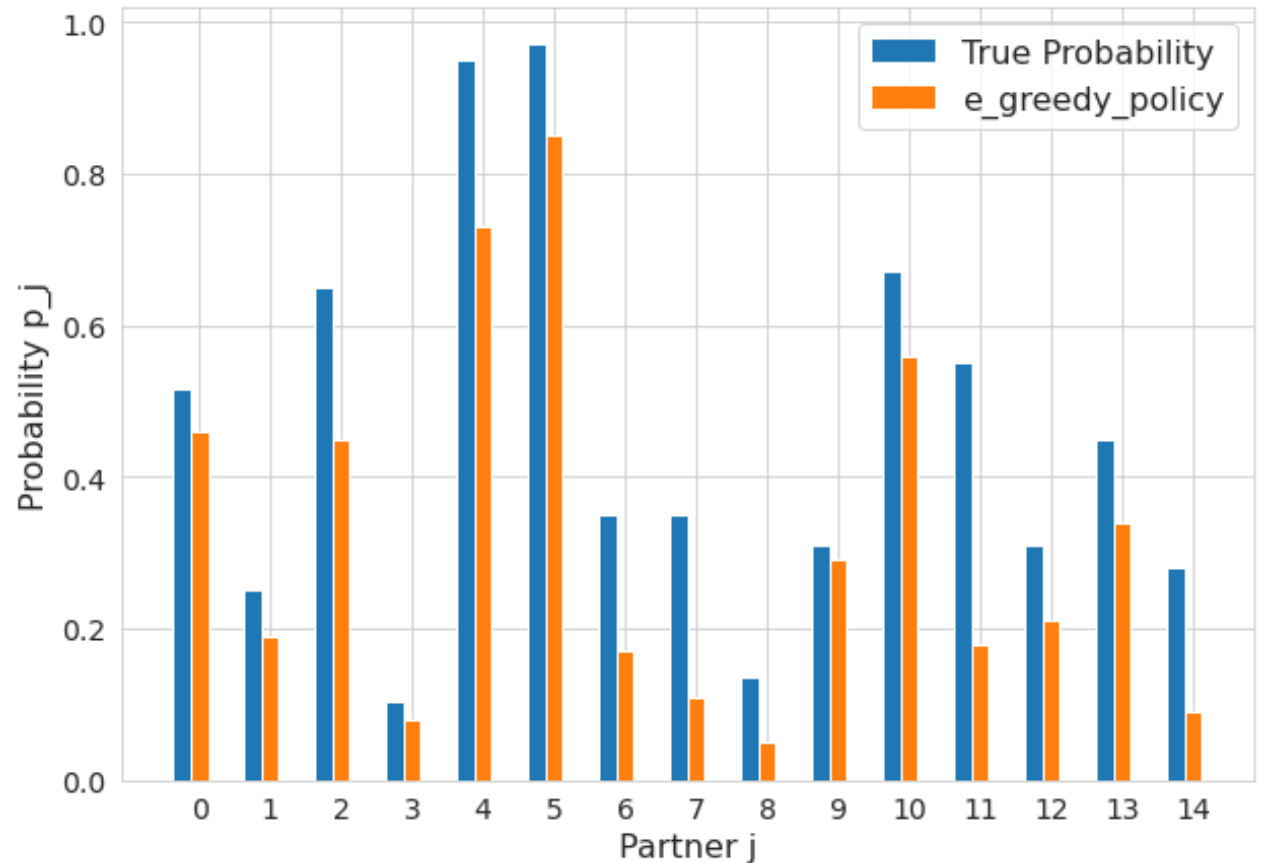


- Solution found on episode 91  
 $- R_{Total} = 0.89$
- However, there is no guarantee to reach always this reward value:

$$\overline{R}_{Total} = 0.67 \quad epi \in [21, 100]$$

# Optimal solution

- We can compute the estimator  $Q_j^*$ 
  - $Q_j^* \approx p_j$
  - Calculate the mean  $Q_j^*$  for the last 10 episodes



## Conclusions and next steps


- We are able to learn the partners with high success probability
- We can use the learning algorithm for pre-selection of partners in ADMM optimization
- Code uploaded in this github [link](#)
- **Next steps:**
  - However, there is still room for improvements...but I'll not be here 😞

Thanks for your attention!



# Multi-Armed Bandit

- Agent learns which arm returns the highest payoff



## **Real-world applications:**

- Clinical trials
- Online Advertising
- Network routing

# Applied RL agent

---

**Algorithm 1:** Applied RL agent

---

Initialize episodes  $e \in \mathbb{E}$ , steps  $n \in \mathbb{N}$ , actions  $a_n \in \mathbb{A}$ , states  $s_n \in \mathbb{S}$ , trading prosumers  $j \in \omega_i$ ;

**for** each episode  $e$  **do**

$E_i^e \leftarrow$  random sample between  $[\underline{E}_i, \overline{E}_i]$ ;

    Initialize step  $n \leftarrow 1$ ;

**while**  $E_i^e \geq s_n$  **do**

        Take action  $a_n = j \in \omega_i$  using policy strategy;

        Observe  $R_n(a_n)$ ,  $s_n$ ;

**if**  $R_n(a_n) = 1$  **then**

$E_{ij}^n \leftarrow$  energy offer by selected prosumer  $j$ ;

**else**

$E_{ij}^n \leftarrow 0$ ;

**end**

$s_n \leftarrow \sum E_{ij}^n$ ,  $R^{total} \leftarrow \sum R_n(a_n)$ ;

        Update the cumulative probability  $\theta(a_n)$ ;

$n \leftarrow n + 1$ ;

**end**

**end**

---

# Total reward performance

<b><i>Algorithm</i></b>	$\overline{R}_{Total}$ <i>epi</i> $\in$ [1, 20]	$\overline{R}_{Total}$ <i>epi</i> $\in$ [21, 100]
<b><i>Random</i></b>	0.49	0.47
<b><i><math>\epsilon</math>-greedy</i></b>	0.50	0.67

# Empirical distribution function across 100 episodes

