


# Reinforcement Learning applications for Energy Analytics and Markets

# Tiago Sousa, Pierre Pinson

## Work done as Postdoc researcher at DTU

# Future Energy markets




— Yes

— No

# Motivation

- We can design Energy markets as Reinforcement Learning




# Outline

- Reinforcement Learning approach
  - Multi-Armed Bandit
- Test case and results
- Conclusions and next steps

# Multi-Armed Bandit

- Agent learns which arm returns the highest payoff



## **Real-world applications:**




- Clinical trials
- Online Advertising
- Network routing

# Multi-Armed Bandit

Agent

participant



|               |  |  |  |
|---------------|---|---|---|
| <i>Step n</i> | <i>Arm 1</i>  | <i>Arm 2</i>  | <i>Arm 3</i>  |
| <b>1</b>      | 1   | 0   | 0   |
| <b>2</b>      | 0   | 1   | 0   |
| <b>3</b>      | 1   | 0   | 0   |
| <b>4</b>      | 0   | 0   | 1   |

actions  $a_n$

Step  $n \neq$  time  $t$

# Multi-Armed Bandit

Agent

participant



Environment

| <i>Step <math>n</math></i> | <i>Arm 1</i> | <i>Arm 2</i> | <i>Arm 3</i> | <i>Reward</i> |
|----------------------------|--------------|--------------|--------------|---------------|
| <b>1</b>                   | 1            | 0            | 0            | 1             |
| <b>2</b>                   | 0            | 1            | 0            | 0             |
| <b>3</b>                   | 1            | 0            | 0            | 0             |
| <b>4</b>                   | 0            | 0            | 1            | 1             |

Total reward

$$R_{Total} = \mathbb{E}(R_n) = \frac{1}{N} \sum R_n$$

$$a_n^* = \operatorname{argmax} R_{Total}$$

# Multi-Armed Bandit

Agent

participant



Environment

| <i>Step <math>n</math></i> | <i>Arm 1</i> | <i>Arm 2</i> | <i>Arm 3</i> | <i>Reward</i> | <i>Energy (kWh)</i> |
|----------------------------|--------------|--------------|--------------|---------------|---------------------|
| <b>1</b>                   | 1            | 0            | 0            | 1             | 10                  |
| <b>2</b>                   | 0            | 1            | 0            | 0             | 0                   |
| <b>3</b>                   | 1            | 0            | 0            | 0             | 0                   |
| <b>4</b>                   | 0            | 0            | 1            | 1             | 5                   |

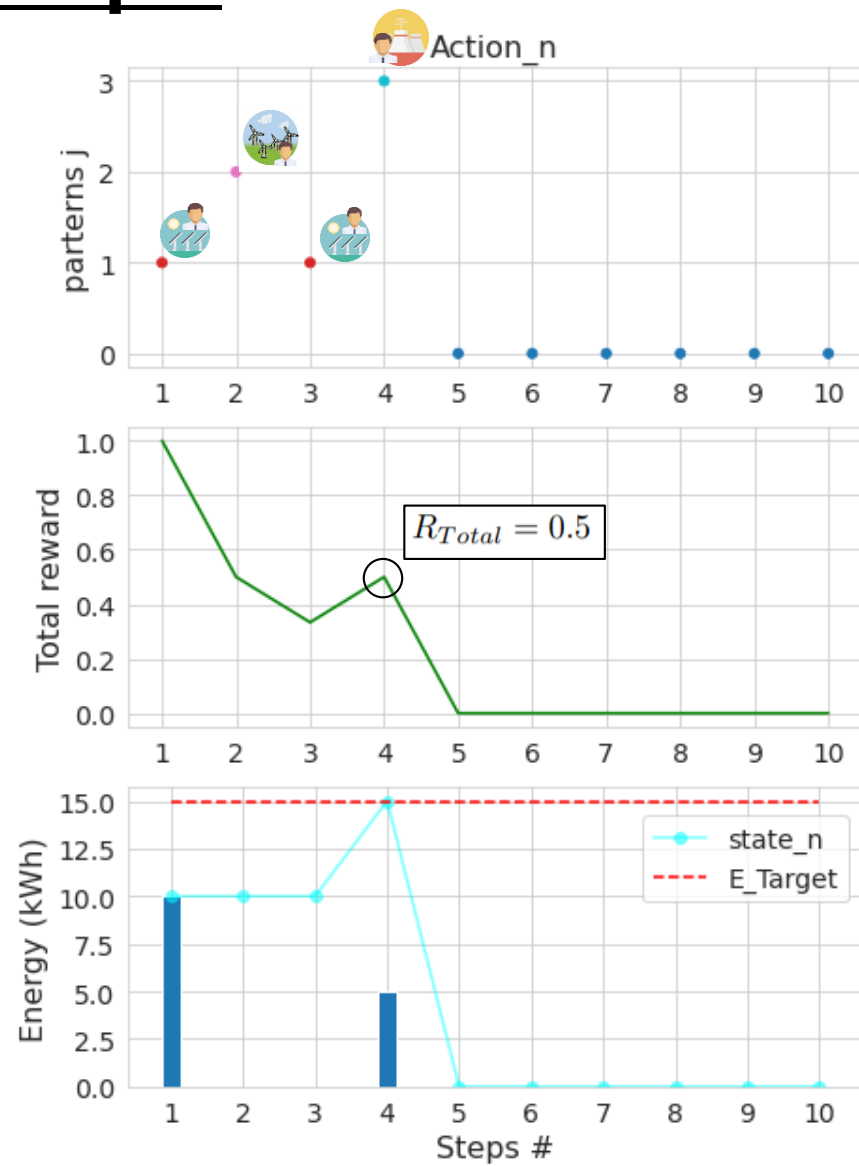
State per step  $n$

$$s_n = \sum E_n(j) R_n$$

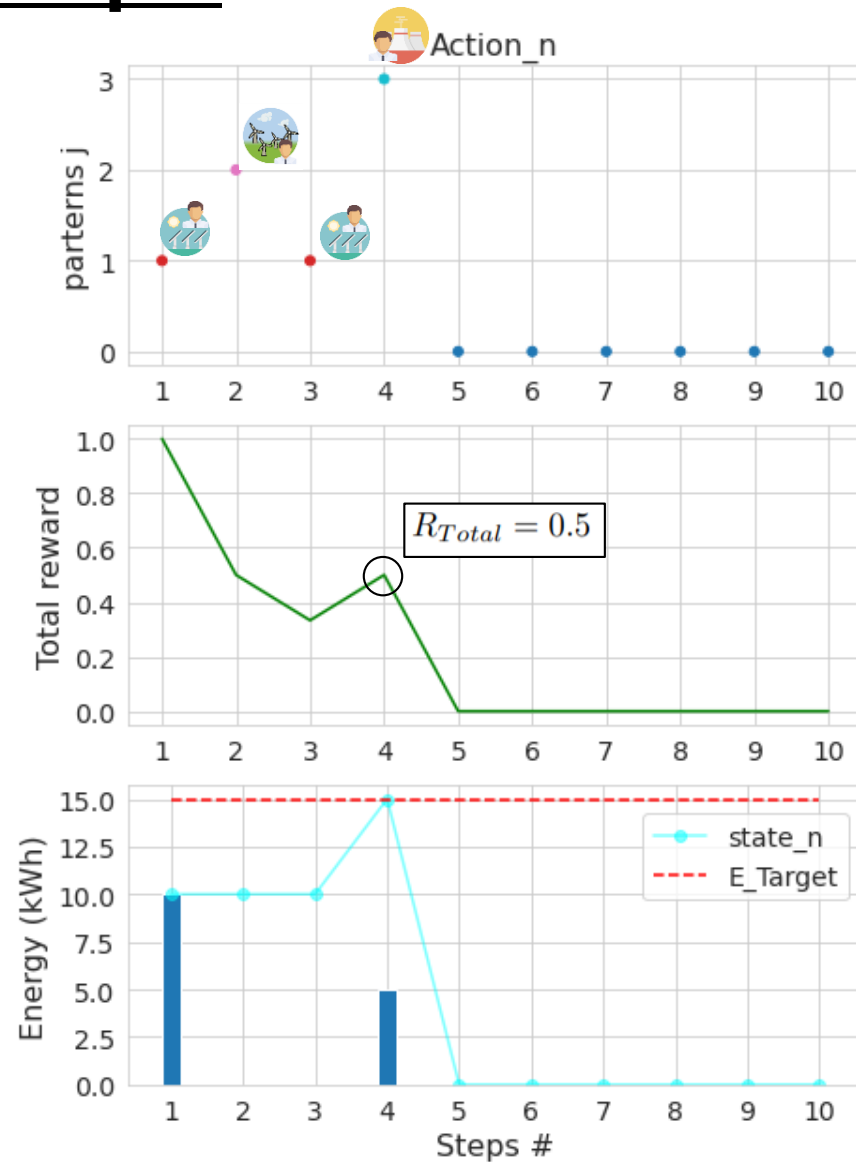
$$s_n \leq E_{Target} \quad \text{Stopping condition}$$



# Example



# Example



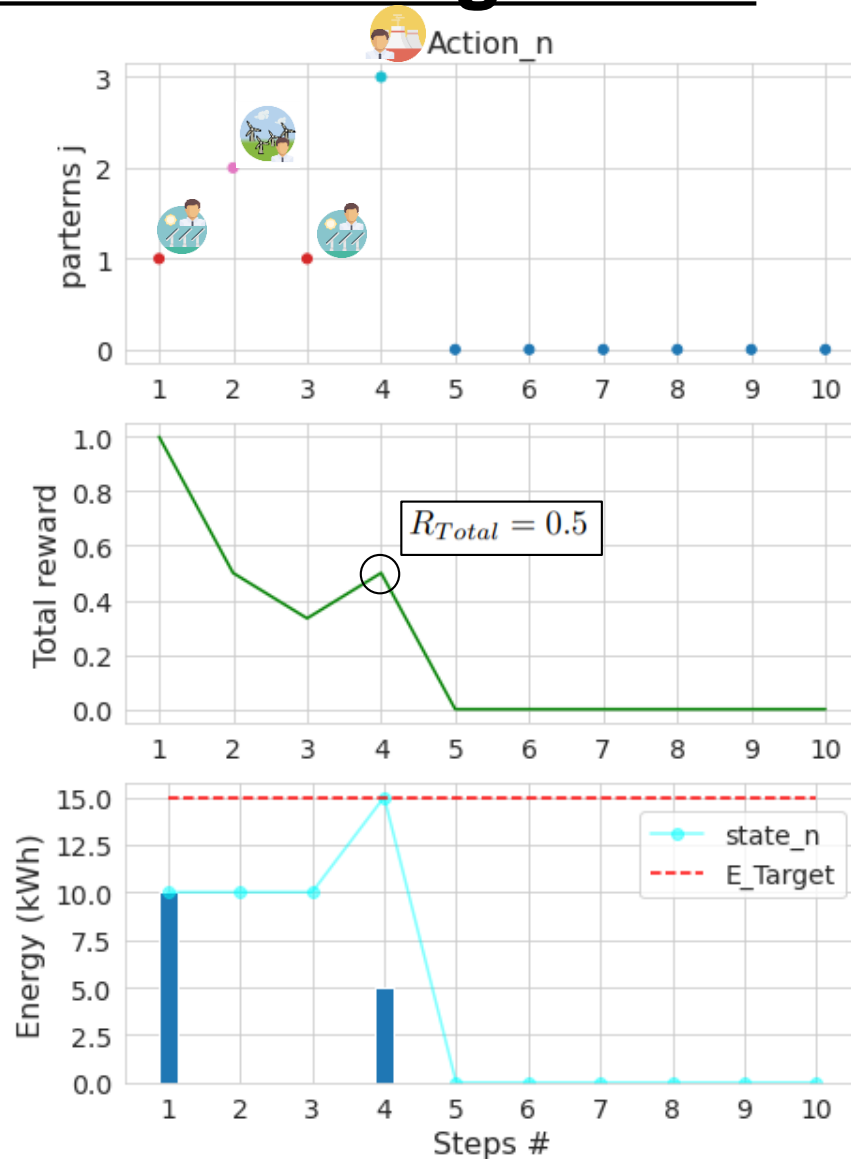
- This iterative process is an episode

- We terminate when

$$s_n = E_{Target}$$

episode = time  $t$

# Translate as Algorithm



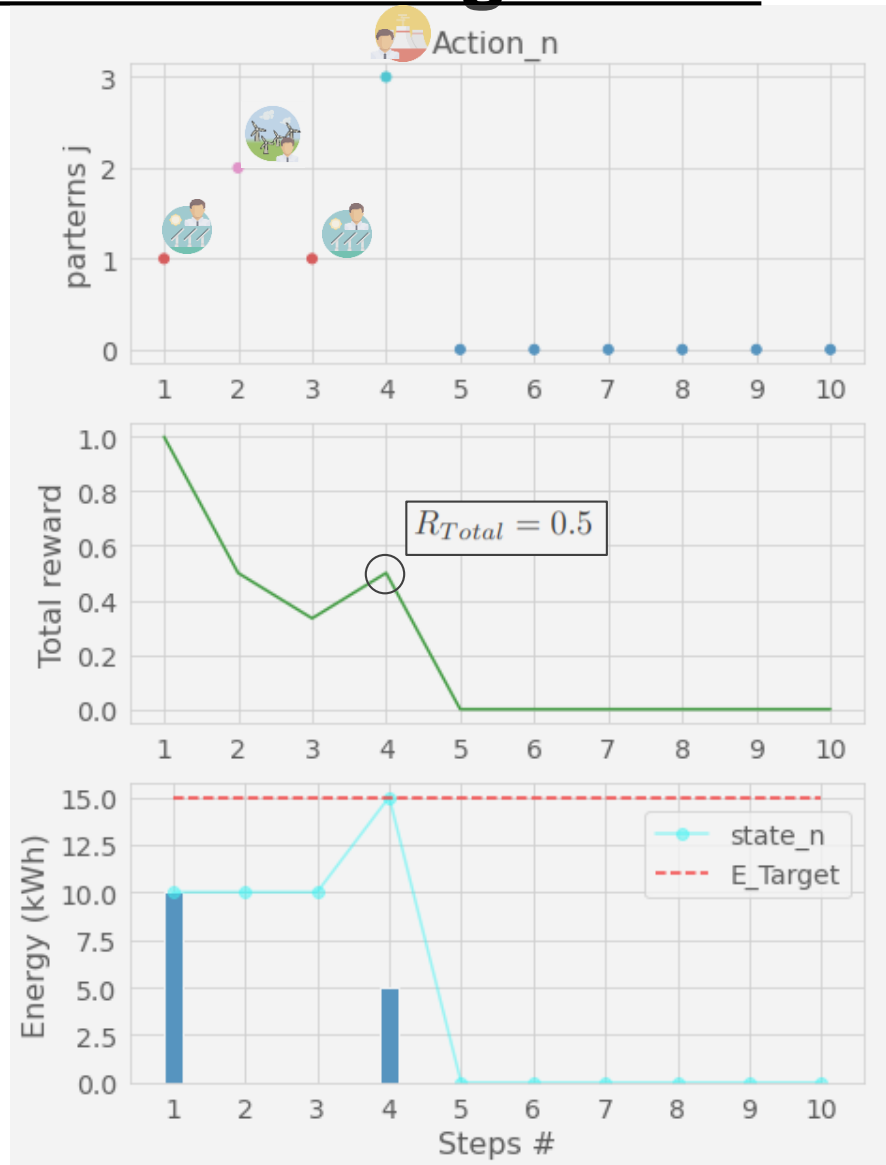
## Algorithm for each episode:

### Algorithm 1: RL Cycle for each episode

```

 $E_{Target} \leftarrow$  random sample from  $[\underline{E}_{Target}, \overline{E}_{Target}]$ ;
Initialize step  $n \leftarrow 1$ ;
while  $s_n \leq E_{Target}$  do
    Take action  $a_n \leftarrow Arm\ j$  (using policy strategy);
    Observe  $R_{Total} \leftarrow \mathbb{E}(R_n)$  ;
    Update  $s_n \leftarrow \sum E_n(a_n)R_n$  ;
     $n \leftarrow n + 1$ ;
end
    
```

# Translate as Algorithm




## Algorithm for each episode:



### Algorithm 1: RL Cycle for each episode

```

 $E_{Target} \leftarrow$  random sample from  $[\underline{E}_{Target}, \overline{E}_{Target}]$ ;
Initialize step  $n \leftarrow 1$ ;
while  $s_n \leq E_{Target}$  do
    Take action  $a_n \leftarrow Arm\ j$  (using policy strategy);
    Observe  $R_{Total} \leftarrow \mathbb{E}(R_n)$  ;
    Update  $s_n \leftarrow \sum E_n(a_n)R_n$  ;
     $n \leftarrow n + 1$ ;
end
    
```



# How to differentiate between episodes?



# Multi-Armed Bandit

- Reward is a random variable

## Bernoulli distribution

$$R_n(j) \sim \text{B}(1, p_j)$$

For a large number of steps  $n$ :

$$R_n(j) \approx p_j$$



| <b>Step <math>n</math></b> | <b>Arm<br/>1</b> |
|----------------------------|------------------|
| <b>1</b>                   | 1                |
| <b>2</b>                   | 0                |
| <b>3</b>                   | 1                |
| <b>4</b>                   | 0                |

## Environment

| <b>Reward</b> |
|---------------|
| 1             |
| -             |
| 0             |
| -             |

# Multi-Armed Bandit

- Reward is a random variable

## Bernoulli distribution

$$R_n(j) \sim \text{B}(1, p_j)$$

For a large number of steps  $n$ :

$$R_n(j) \approx p_j$$



| Step $n$ | Arm 1 | Reward |
|----------|-------|--------|
| 1        | 1     | 1      |
| 2        | 0     | -      |
| 3        | 1     | 0      |
| 4        | 0     | -      |

## Environment

## Action-Value function

Estimator of  $p_j$  for every arm  $j$

$$Q_n(j) = \frac{1}{N_j} \sum R_n(j) = \hat{p}_n(j)$$

# Mathematical Formulation

---

**Algorithm 2:** Complete algorithm

---

Initialize episodes  $e \in \mathbb{E}$ , steps  $n \in \mathbb{N}$ , actions  $a_n \in \text{arms } \mathbb{J}$  ;

**for** *each* episode  $e$  **do**

$E_{Target} \leftarrow$  random sample from  $[\underline{E}_{Target}, \overline{E}_{Target}]$ ;

    Initialize step  $n \leftarrow 1$ ;

**while**  $s_n \leq E_{Target}$  **do**

        Take action  $a_n \leftarrow \text{Arm } j$  (using policy strategy);

        Observe  $R_{Total} \leftarrow \mathbb{E}(R_n)$  ;

        Update  $s_n \leftarrow \sum E_n(a_n)R_n$  ;

$n \leftarrow n + 1$ ;

        Update every  $Q_n(j) \leftarrow \mathbb{E}(R_n(j)) = \hat{p}_n(j)$  ;

**end**

    Propagate to the next episode  $Q^{e+1}(j) \leftarrow \mathbb{E}(Q^e(j))$  ;

**end**

---



# Mathematical Formulation

---

**Algorithm 2:** Complete algorithm

---

Initialize episodes  $e \in \mathbb{E}$ , steps  $n \in \mathbb{N}$ , actions  $a_n \in \text{arms } \mathbb{J}$  ;

**for** *each episode*  $e$  **do**

$E_{Target} \leftarrow$  random sample from  $[\underline{E}_{Target}, \overline{E}_{Target}]$ ;

    Initialize step  $n \leftarrow 1$ ;

**while**  $s_n \leq E_{Target}$  **do**

        Take action  $a_n \leftarrow \text{Arm } j$  (using policy strategy);

        Observe  $R_{Total} \leftarrow \mathbb{E}(R_n)$  ;

        Update  $s_n \leftarrow \sum E_n(a_n)R_n$  ;

$n \leftarrow n + 1$ ;

        Update every  $Q_n(j) \leftarrow \mathbb{E}(R_n(j)) = \hat{p}_n(j)$  ;

**end**

    Propagate to the next episode  $Q^{e+1}(j) \leftarrow \mathbb{E}(Q^e(j))$  ;

**end**


---

## Experience replay as NN

- We start with *batch* of episodes with **no propagation**
- After, we compute the average Action-value for each arm  $j$


## More to say!!!

- How to adopt other propagation strategies between episodes
- Learning algorithms estimate the Action-Value function  $Q_n(j)$ 
  - Random
  - $\epsilon$ -greedy
  - Thompson Sampler
  - Upper Confidence Bound




# Test case

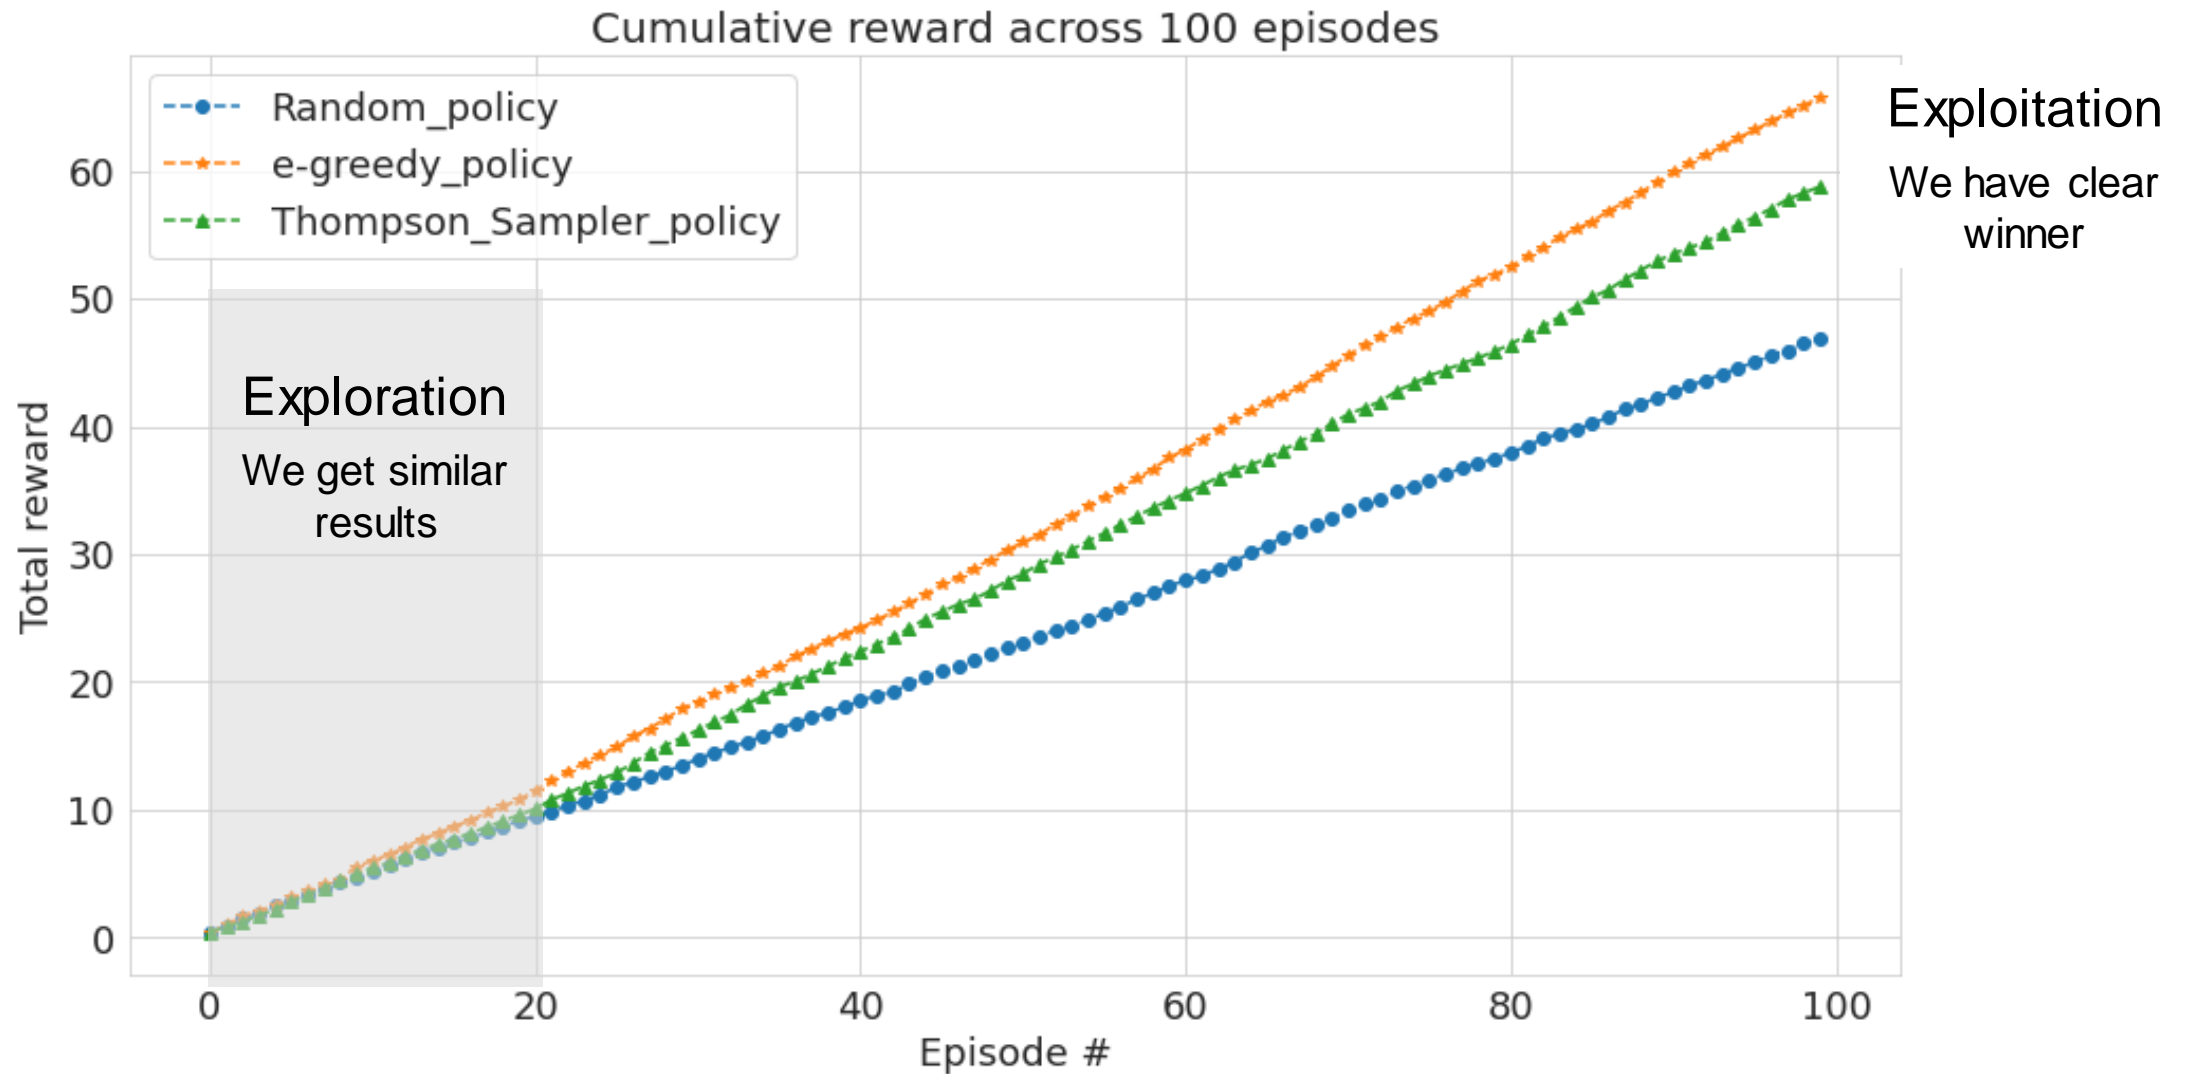
- Case with 15 partners  $j$
- We used 100 episodes
  - $E_{Target} = 15$  kWh
- Learning algorithms:
  - Random
  - $\epsilon$ -greedy
  - Thompson Sampler



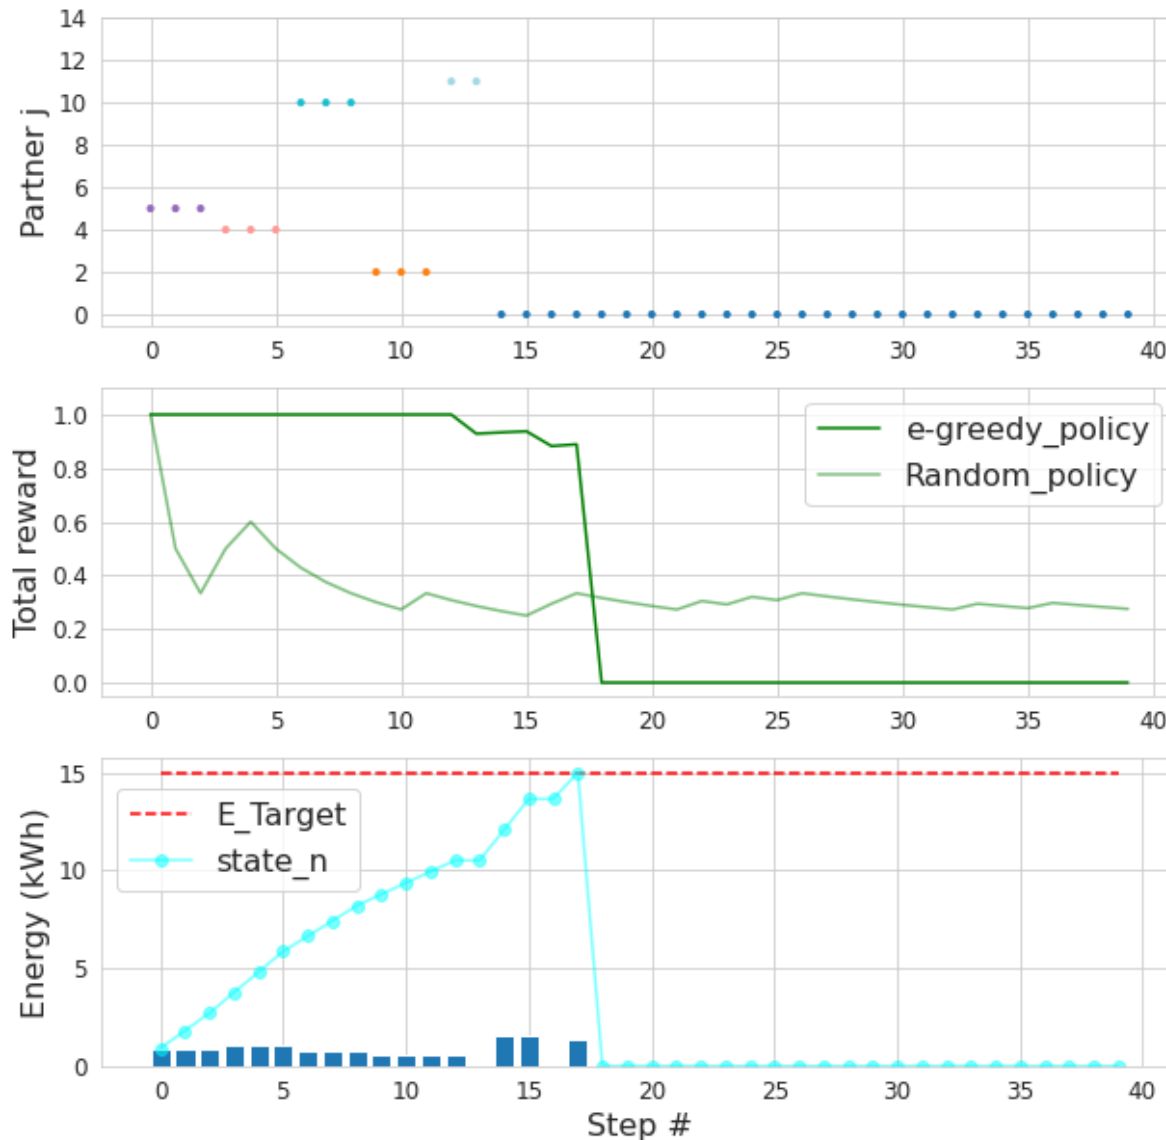
# Test case



# Compare strategies



# Epsilon-Greedy algorithm



- Solution found on episode 91

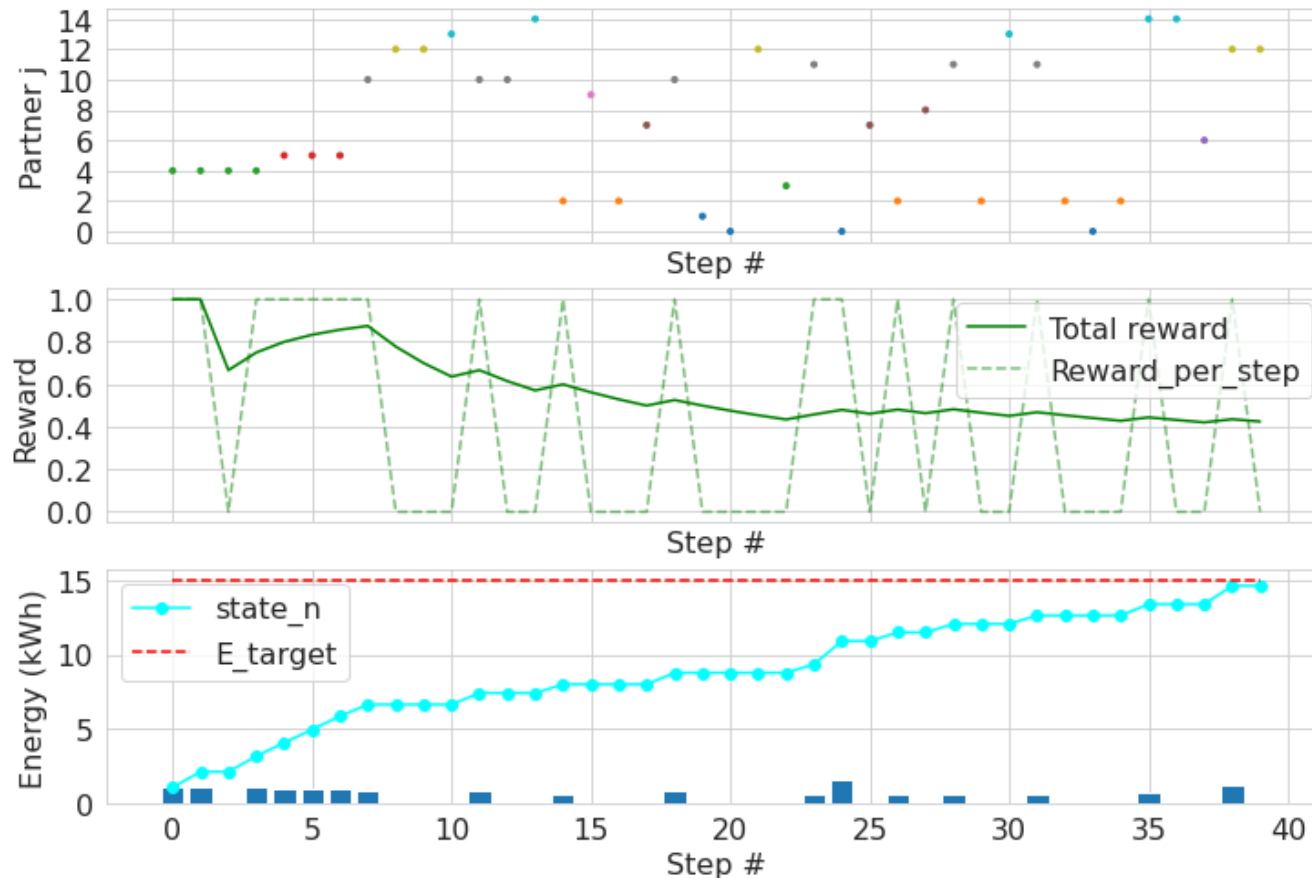
$$- R_{Total} = 0.89$$

- However, there is **no guarantee** to reach always this reward value:

$$\overline{R}_{Total} = 0.67 \quad epi \in [21, 100]$$

# Thompson Sampler algorithm

- We can also have ‘bad’ results even in the **exploitation phase**

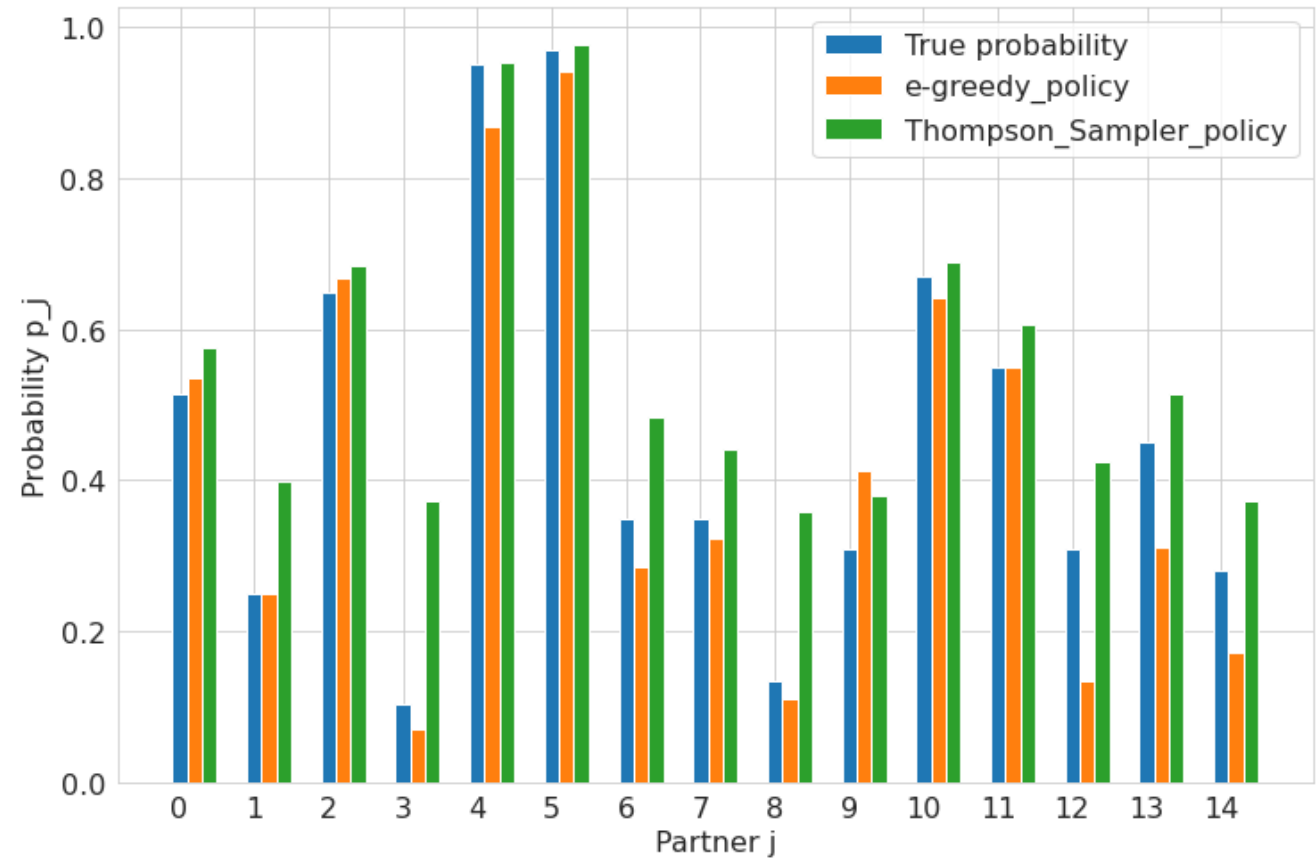


- Solution found on episode 91

$$R_{Total} = 0.41$$

# Optimal solution

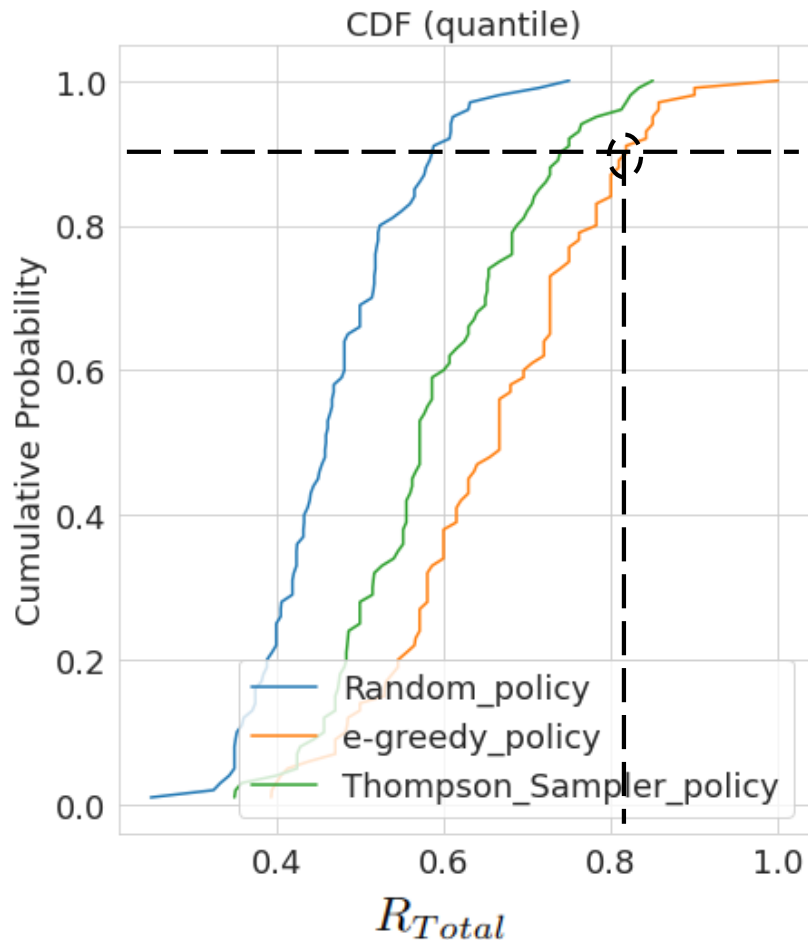
- We can compute the estimator  $Q_j^*$ 
  - $Q_j^* \approx p_j$
  - Calculate the mean  $Q_j^*$  for the last 10 episodes





# Alternative approach

- In fact, we can retrieve the CDF of the  $R_{Total}$  for the 100 episodes



## 90% Percentile

- Filter the best episodes
- We can define an upper bound:

$$R_{Total} \geq 0.82$$



- We would then compute the estimator:
  - $Q_j^* \approx p_j$

## Conclusions and next steps

- We are able to learn the partners with high success probability  $Q_j^* \approx p_j$
- Easy way to build a learning agent via the Q-value (Action-value) functions
- Code available in my GitHub repo ([link here](#))
- **Next steps:**
  - Estimate the  $Q_j^*$  using the CDF of the  $R_{Total}$
  - Improve the Experience replay for the propagation
  - Assess the performance via a validation phase

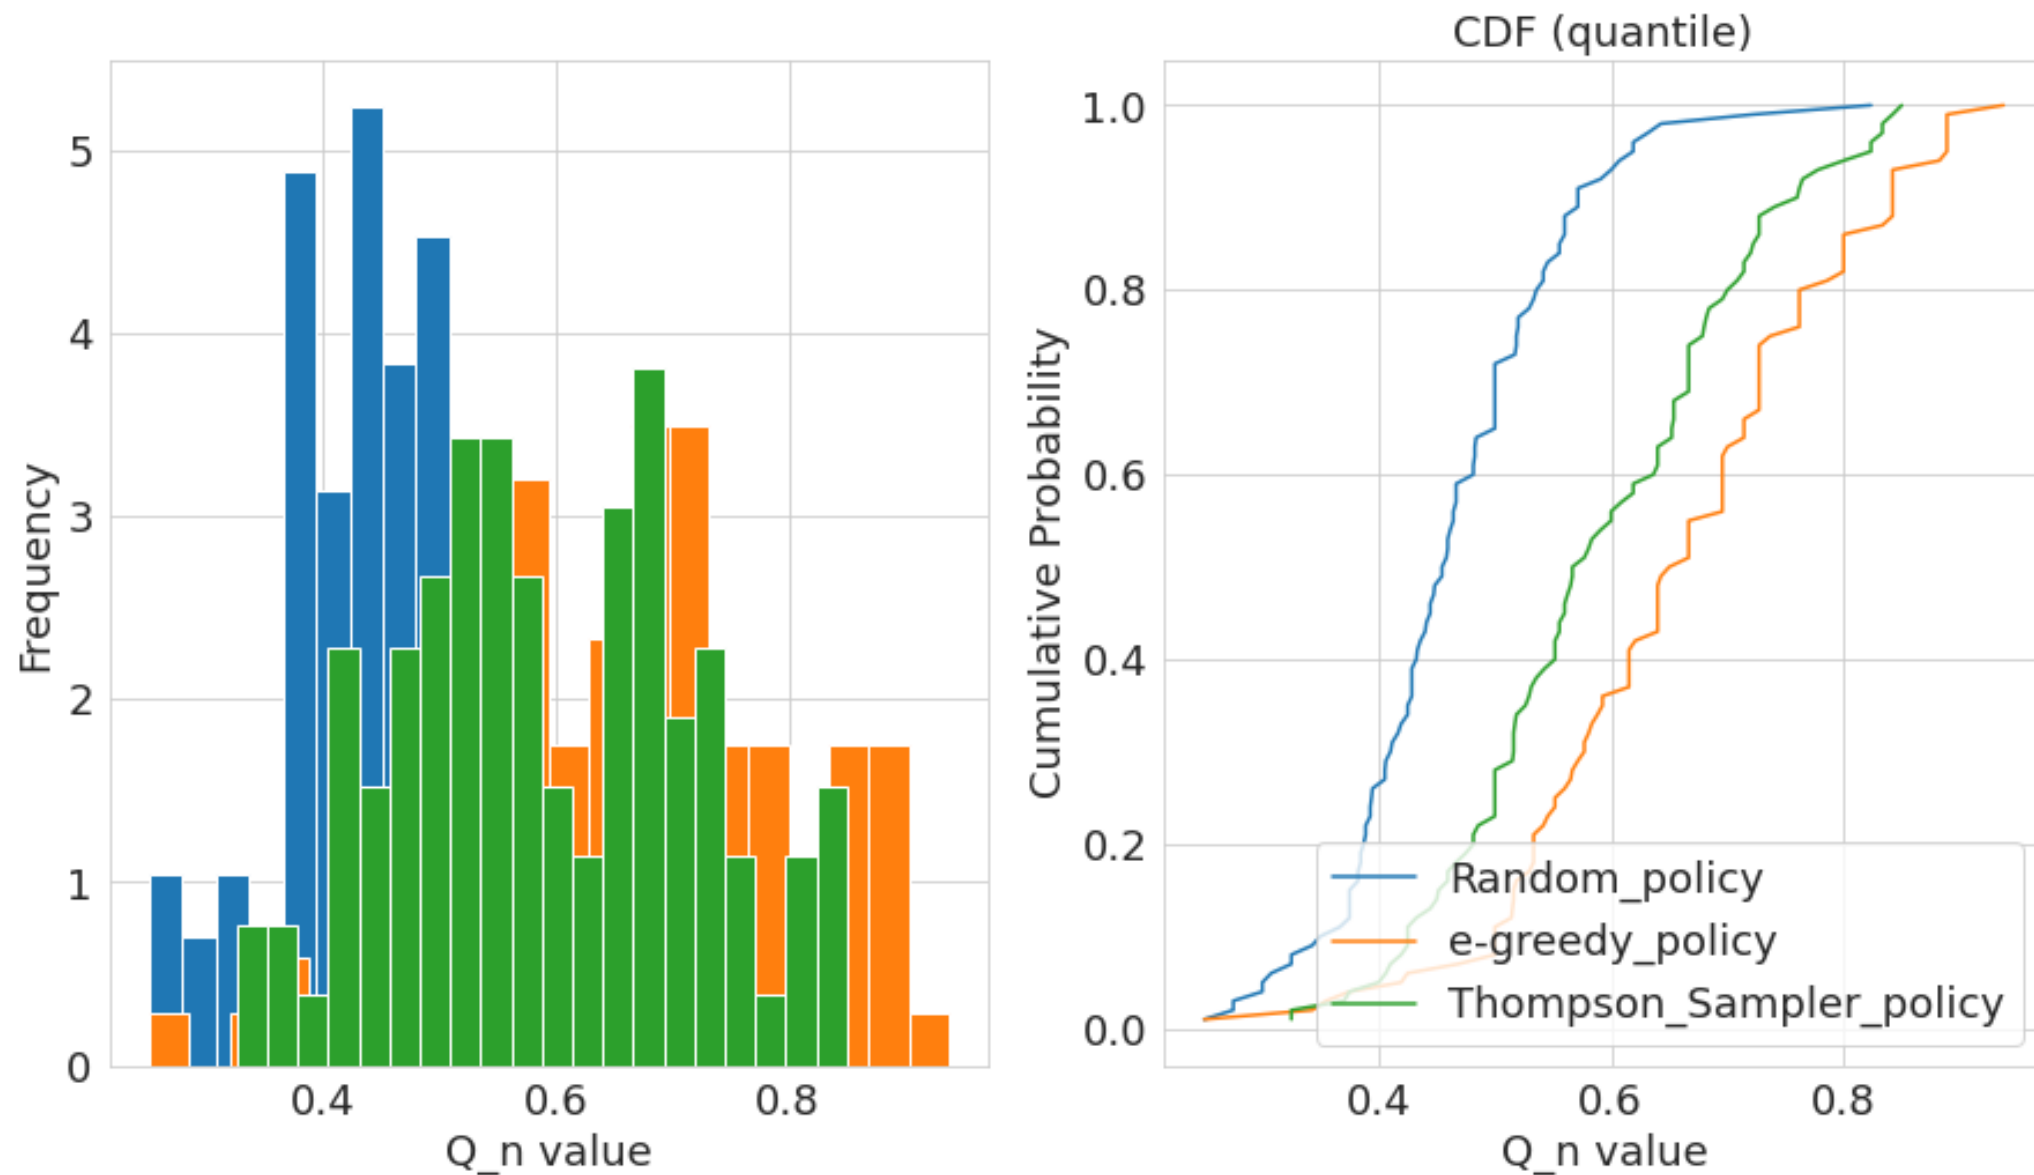
**Thanks for your attention!**





# Total reward performance

| <i><b>Algorithm</b></i>                    | $\overline{R}_{Total}$<br><i><math>epi \in [1, 20]</math></i> | $\overline{R}_{Total}$<br><i><math>epi \in [21, 100]</math></i> |
|--|---|---|
| <i><b>Random</b></i>                       | 0.49  | 0.47  |
| <i><b><math>\epsilon</math>-greedy</b></i> | 0.50  | 0.67  |

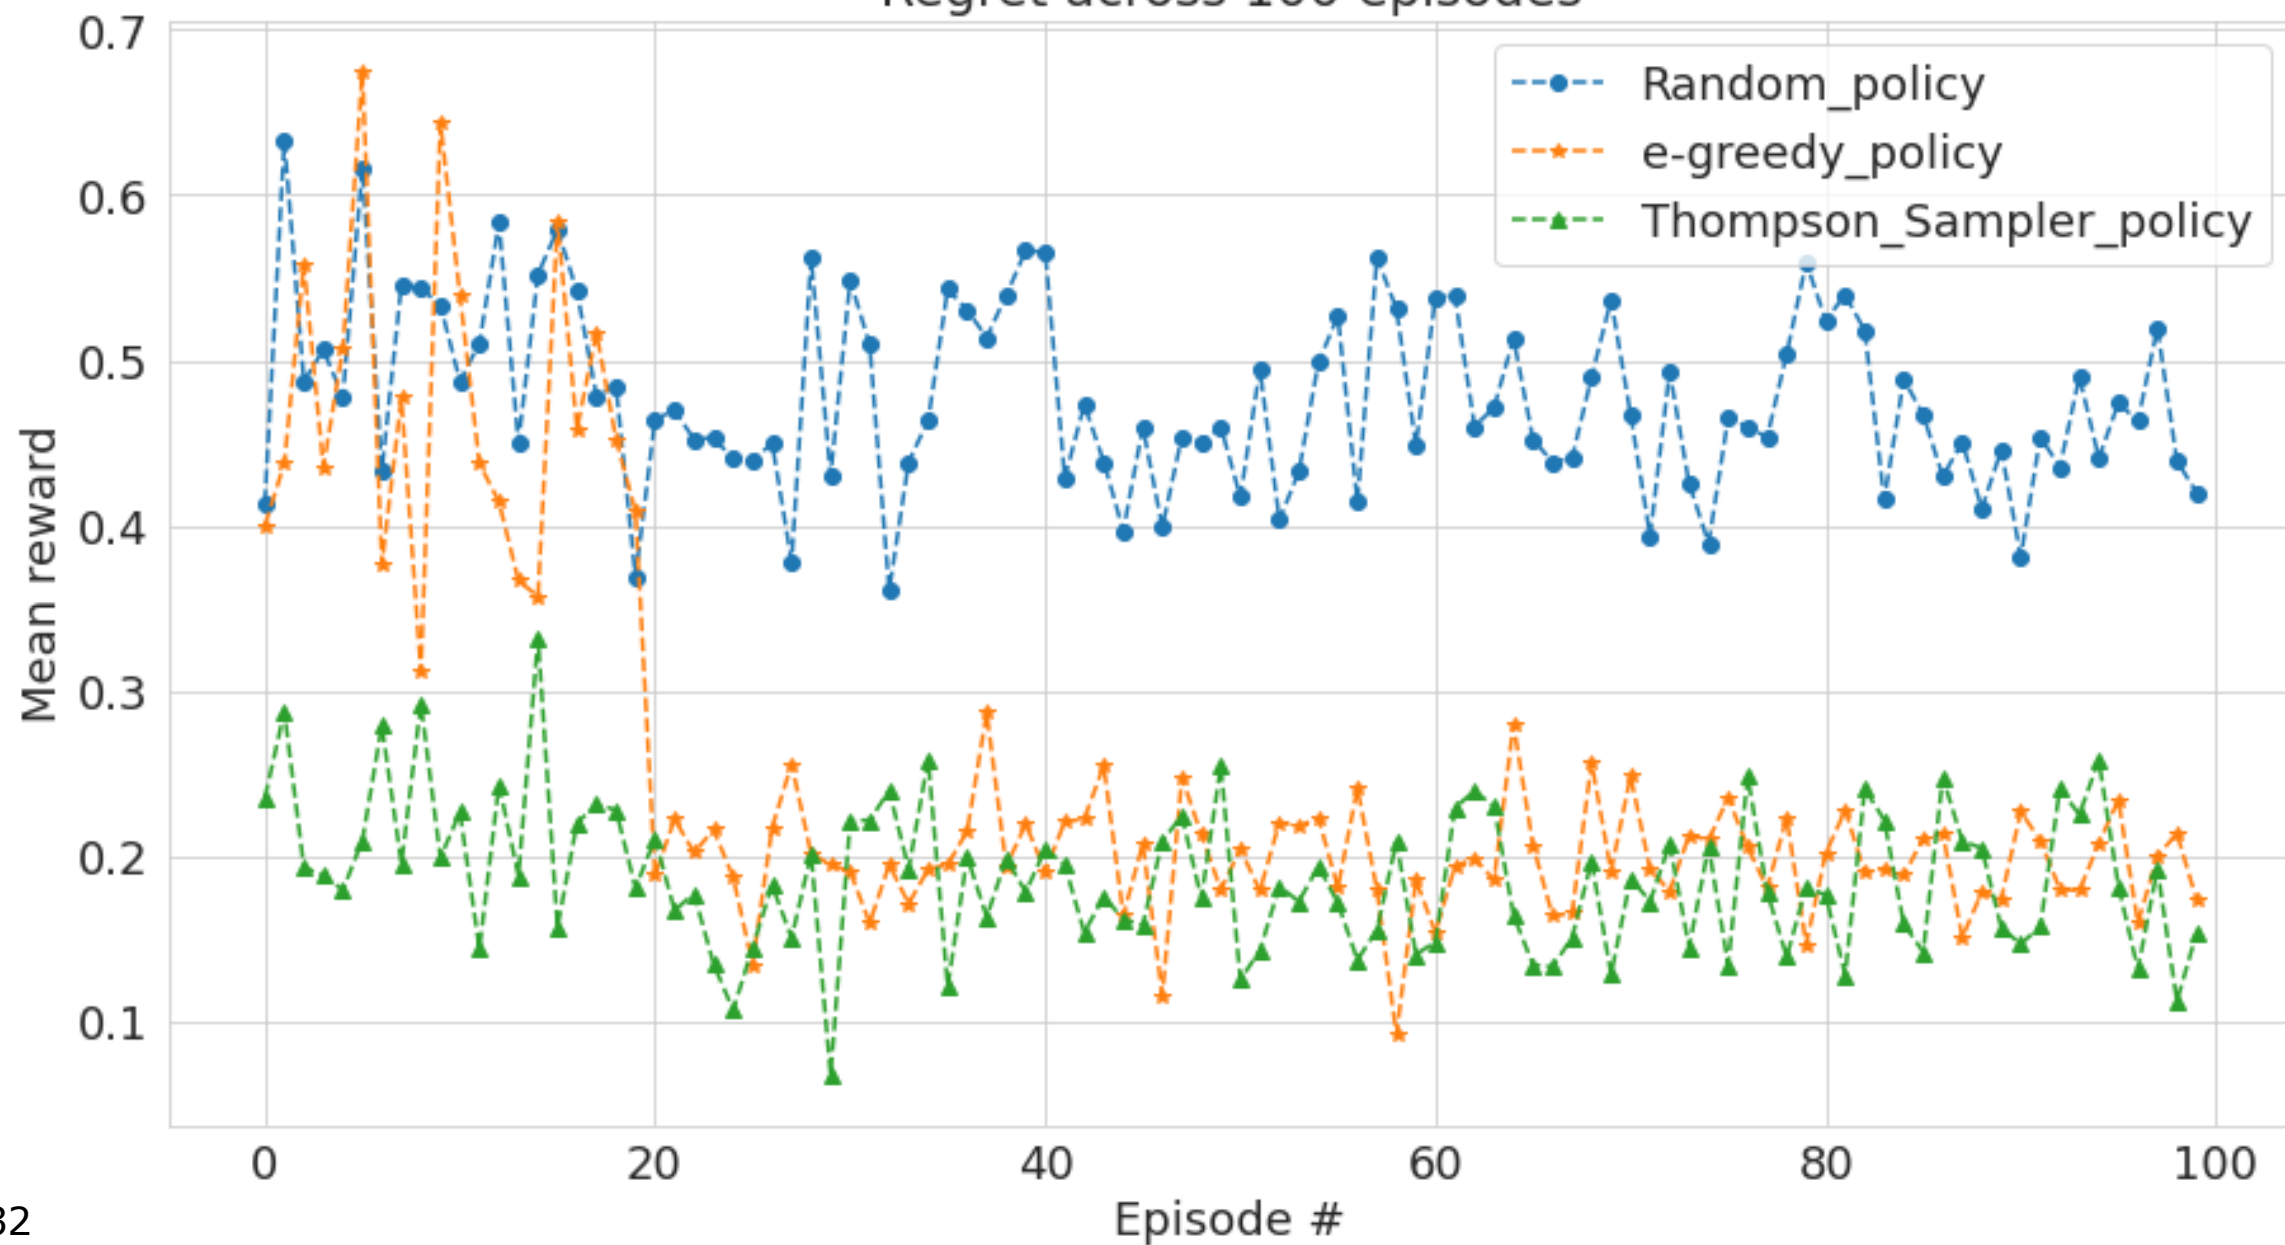
# Empirical distribution function across 100 episodes







Regret across 100 episodes





Optimal estimator  $Q(\text{arm}_j)$  per RL\_agent